

BOARD BRIEF — V2, RED-TEAM HARDENED

Wallace Ad Intelligence Engine

A statistical decision engine that finds advertising candidates, scores them against measured evidence, survives adversarial review, and briefs a human — who alone approves, launches, and spends.

NO AUTO-SPEND

EVIDENCE-LED

7 RED-TEAM ROUNDS SURVIVED

WHAT CHANGED SINCE V1

V1 recommended ads. V2 has to prove them.

Every recommendation now carries its own audit trail: a permanent campaign registry, an immutable pre-launch baseline, a 0–100 opportunity score built from six measured components, dual confidence grades — and it all survived seven rounds of adversarial review by an independent red-team model.

6

Score components

Conversion propensity, demand momentum, position gap, audience, readiness, strategic fit.

2

Confidence grades

Evidence Confidence (A–D) and Channel Readiness graded separately — uncertainty can never earn points.

7

Red-team rounds

Each round found real inflation. Each fix made scores drop. That is the point.

\$0

Spent so far

Nothing launches or spends without a signed human preflight.

CORE RULE — Missing data lowers confidence. It never earns neutral points.

Six stages. One permanent record. One human gate.

1 — FIND

Scouts scan GA4 session-level event rates, GSC 28-day windows, listing-demand rollups, and SEMrush competitor ads for candidate pages.

2 — VET

Seven hard gates: page health, evergreen destination, destination class, data minimums, duplication, budget, history. Failures keep their reason forever.

3 — SCORE

0–100 opportunity score from six evidence components, shrunk toward site baselines, with Wilson lower-bound guards on small samples.

4 — BRIEF

Weekly brief with execution kit (copy, CTA, image direction, exact UTM link), budget, and both confidence grades — committed to the private repo.

5 — APPROVE & LAUNCH

Human-only. Fair-housing preflight, Adwerx checkout economics, and the native Adwerx campaign ID are recorded before a dollar moves.

6 — TRACK & LEARN

Friday tracker joins spend, delivery, and GA4 attribution on exact IDs only, and issues day-7/30/60/90 verdicts against an immutable baseline.

Each source has one job — and one blind spot we name.

SOURCE	JOB	WHAT IT CAN NOT PROVE (AND THE ENGINE SAYS SO)
GA4	Session-scoped event rates for five qualified actions; landing-page attribution	Event overlap within a session — handled with QAV lower/upper bounds
Search Console	Impressions, clicks, position; 28-day current vs prior windows	Rising clicks $\neq$ rising demand — classified separately (see demand flag)
SEMrush MCP	Competitor ad pressure, keyword context	Competitor ads prove competitor spending — not that our ad will work
Adwerx	The only source of real reach, impressions, price — captured at checkout	Nothing is assumed before checkout; organic traffic is only a capped proxy
n8n	Two workflows: weekly recommender + Friday tracker (dry-run, versioned)	—
GitHub (private)	Registry, immutable baselines, briefs, scorecards, rejection history	—
Twilio SMS	Outbound operational alerts only	Inbound SMS commands disabled (unauthenticated channel)

QAV: five weighted actions, honest about overlap

Raw "leads" overcount: one visitor can fire multiple events. QAV weights each qualified action (contact form 1.0, listing conversion 1.0, custom form 0.6, registration 0.35, click-to-call 0.25) and reports a **bounded range**: the lower bound assumes maximum overlap (each session credited only its most valuable action), the upper bound assumes none. Scores use the lower bound. A wide range cuts confidence.

$$QAV_{low} \leq \text{true value} \leq QAV_{high} \text{ — scored on } QAV_{low}$$

Proven by unit test

Synthetic page: 100 sessions, 10 event-sessions of each of the 5 types. Engine returns $QAV_{low} = 10.0$ (all overlapping), $QAV_{high} = 32.0$ (all separate). The red-team predicted the original code would return 32 as the "lower bound" — it did. Fixed, tested, committed.

Why the board should care

Before this fix, the flagship "luxury" candidate scored 80/A. Its true lower-bound evidence supports 39/D. That is a \$500/month decision made honest by arithmetic.

THE OPPORTUNITY SCORE

100 points, none of them free

COMPONENT	MAX	GUARDRAIL THAT STOPS INFLATION
Conversion propensity (3 tiers)	30	Zero observed events = zero points. One event caps at $\frac{1}{3}$ of its tier; Wilson 80% lower bound must beat the site baseline to score at all.
Demand & momentum	20	Absolute click caps (<20 clicks → max 5). Clicks up while impressions fall = VISIBILITY-ONLY, hard 5-point ceiling.
Position gap	15	Pages already ranking 1–3 get 0 — paid media adds least where organic already wins. Sweet spot is positions 5–10.
Addressable audience	15	Organic traffic is only a PROXY, capped at 5/15. Full points require an actual Adwerx reach forecast from checkout.
Readiness	10	Page loads, tracking fires, forms verified.
Strategic fit	10	Brokerage-owned, evergreen destinations only.

Funding needs **both** high opportunity and high confidence. High score + low confidence = controlled test at most. Low + low = "needs data," never "maybe."

Seven rounds with a hostile reviewer — every score fell

An independent red-team model attacked the engine's math for seven rounds. Each round it found a real way uncertainty was being converted into points. Falling scores are the proof the system works: the engine now refuses to spend on evidence it cannot defend.

CANDIDATE	INFLATED SCORE	CORRECTED	WHAT WAS INFLATING IT
Luxury properties	80 / A	39 / D	Overlap-blind QAV; visibility-only click growth counted as demand
Commercial resources	58 / D	35 / D	One single contact form saturated an 18-point tier; unverified audience points
Louisville / Oneida	69 / C	33 / 30 / C	Zero measured conversions — all propensity points came from the statistical prior; max points for ranking #3 already

RESULT — The engine currently recommends funding NOTHING. That is the honest answer.

Sitewide demand softening — seasonality unresolved

Total site search impressions fell 755,200 → 710,279 (−6%) over 28 days, with 351 of 433 tracked pages declining. Until year-over-year data resolves whether this is seasonal, every page's "momentum" is classified against the site trend:

SIGNAL	CLASSIFICATION	SCORE TREATMENT
Impressions rising absolutely and vs site	TRUE DEMAND GROWTH	Full momentum points (up to 20)
Impressions declining, but less than the site	RELATIVE RESILIENCE	Small contextual credit (≤5) — never labeled rising demand
Clicks rising while impressions fall	VISIBILITY-ONLY	Hard 5-point ceiling — CTR/rank improvement, not market demand

This banner now heads every weekly brief so no one mistakes a resilient page for a growing market.

Money is reserved, aged, and released — never assumed

\$1,500

Consumer cap /mo

Max three concurrent campaigns. \$500 careers lane held separately with its own rules.

7 / 14

Proposal lifecycle

Proposals reserve budget for 7 days and are revalidated weekly; approvals go stale after 14 days unlaunched. Superseded ideas release their money automatically.

\$500

Currently held back

Unspent budget stays unspent. The engine's own brief says: evidence doesn't justify forcing it out the door.

- ✓ Every proposal is a registry record with a permanent WAL ID, exact UTM, and (at launch) the native Adwerx campaign ID.
- ✓ An immutable pre-launch baseline file is committed before approval — day-30/60/90 verdicts compare against that frozen snapshot, never a re-computed one.
- ✓ Spend reports that cannot be matched to a registered campaign by exact ID are flagged UNVERIFIED SPEND and excluded from all verdicts until a human resolves them.

An operational pilot, honestly labeled as one

The model recommends funding nothing. Leadership still needs to prove the pipes: delivery, attribution, spend reconciliation, and verdicts. So Pilot 1 is encoded as a **manual operational-validation override** — not a model pick, and its outcome can never be used to recalibrate the model as if it were one.

The record (WAL-2026-007, commercial)

```
campaignMode: operational_validation
modelRecommendation: WATCH — preserved forever
manualOverride: true • maxBudget: $500
excludedFromScoreCalibration: true
```

The decision sequence

- 1 — Inspect Adwerx checkout (no purchase) for commercial and the luxury fallback: real price, promised impressions, channels, audience, frequency.
- 2 — Choose commercial unless its economics are materially worse.
- 3 — One campaign. \$500 hard cap. Signed fair-housing preflight first.

Success = plumbing, not ROI: native ID reconciles, impressions reconcile, exact UTM lands in GA4, sessions attach only to this pilot, snapshots don't duplicate, and the day-7 verdict correctly separates delivery vs tracking vs engagement failures.

Every Friday, every campaign gets graded — or honestly ungraded

Day 7

Operational chain: serving?
tracked? attributed?
Distinguishes delivery, tracking,
and engagement failures.

Day 30

First performance read vs the
immutable baseline and the
launch goal.

Day 60

Keep, modify, scale, hold, or
stop — with the reason
recorded.

Day 90

Final score. The lesson feeds
the next vetting run's history
gate.

- ✓ Joins are exact-only: native Adwerx ID, WAL registry ID, or exact utm_campaign. No fuzzy matching, ever.
- ✓ Cumulative delivery reports replace prior snapshots — never blindly summed.
- ✓ "Insufficient data" is a legitimate verdict. The tracker will say it rather than guess.
- ✓ Failure alerts go out by SMS; weekly results are committed to the private repo.

The engine advises. People decide. Always.

1 — ENGINE

Finds, vets, scores, briefs. Dry-run only.

2 — HOWARD / KATE

Preflight checklist: fair housing, branding, checkout economics, exact UTM.

3 — ADWERX

Launch happens by hand, in the platform, after sign-off.

4 — ENGINE

Tracks, reconciles, grades, and remembers.

- ✓ No API path exists for the engine to create, modify, or fund a live campaign.
- ✓ Fair-housing review is a named human sign-off on every campaign — no protected-class targeting or proxies.
- ✓ Inbound SMS control is disabled; the registry is protected against concurrent-write corruption.

BOTTOM LINE

The engine earned the right to be trusted by refusing to flatter its own ideas.

Seven adversarial rounds, every inflation removed, zero dollars spent, a full audit trail — and one honestly-labeled \$500 operational pilot waiting on a single human step.

NEXT MOVE — Inspect Adwerx checkout (no purchase), record the economics, sign the preflight.